

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as seaborn
import warnings
warnings.filterwarnings("ignore")
```

```
df=pd.read_excel("Churn_Modelling.xlsx")
```

```
df.head()
```

| | RowNumber | CustomerId | Surname | CreditScore | Geography | Gender | Age | Tenure | |
|---|-----------|------------|----------|-------------|-----------|--------|-----|--------|----|
| 0 | 1 | 15634602 | Hargrave | 619 | France | Female | 42 | 2 | |
| 1 | 2 | 15647311 | Hill | 608 | Spain | Female | 41 | 1 | 8 |
| 2 | 3 | 15619304 | Onio | 502 | France | Female | 42 | 8 | 15 |
| 3 | 4 | 15701354 | Boni | 699 | France | Female | 39 | 1 | |
| 4 | 5 | 15737888 | Mitchell | 850 | Spain | Female | 43 | 2 | 12 |

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber              10000 non-null  int64
1   CustomerId             10000 non-null  int64
2   Surname                10000 non-null  object
3   CreditScore            10000 non-null  int64
4   Geography              10000 non-null  object
5   Gender                 10000 non-null  object
6   Age                   10000 non-null  int64
7   Tenure                 10000 non-null  int64
8   Balance                10000 non-null  float64
9   NumOfProducts          10000 non-null  int64
10  HasCrCard              10000 non-null  int64
11  IsActiveMember         10000 non-null  int64
12  EstimatedSalary        10000 non-null  float64
13  Exited                 10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
df.drop(columns=["RowNumber","CustomerId","Surname"],inplace=True)
```

```
df
```

| | CreditScore | Geography | Gender | Age | Tenure | Balance | NumOfProducts | HasCrC |
|------|-------------|-----------|--------|-----|--------|-----------|---------------|--------|
| 0 | 619 | France | Female | 42 | 2 | 0.00 | 1 | |
| 1 | 608 | Spain | Female | 41 | 1 | 83807.86 | 1 | |
| 2 | 502 | France | Female | 42 | 8 | 159660.80 | 3 | |
| 3 | 699 | France | Female | 39 | 1 | 0.00 | 2 | |
| 4 | 850 | Spain | Female | 43 | 2 | 125510.82 | 1 | |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 9995 | 771 | France | Male | 39 | 5 | 0.00 | 2 | |
| 9996 | 516 | France | Male | 35 | 10 | 57369.61 | 1 | |
| 9997 | 709 | France | Female | 36 | 7 | 0.00 | 1 | |
| 9998 | 772 | Germany | Male | 42 | 3 | 75075.31 | 2 | |
| 9999 | 792 | France | Female | 28 | 4 | 130142.79 | 1 | |

10000 rows × 11 columns

```
from sklearn.preprocessing import LabelEncoder
```

```
le=LabelEncoder()  
df["Geography"]=le.fit_transform(df["Geography"])
```

```
le1=LabelEncoder()  
df["Gender"]=le1.fit_transform(df["Gender"])
```

df

| | CreditScore | Geography | Gender | Age | Tenure | Balance | NumOfProducts | HasCrC |
|------|-------------|-----------|--------|-----|--------|-----------|---------------|--------|
| 0 | 619 | 0 | 0 | 42 | 2 | 0.00 | 1 | |
| 1 | 608 | 2 | 0 | 41 | 1 | 83807.86 | 1 | |
| 2 | 502 | 0 | 0 | 42 | 8 | 159660.80 | 3 | |
| 3 | 699 | 0 | 0 | 39 | 1 | 0.00 | 2 | |
| 4 | 850 | 2 | 0 | 43 | 2 | 125510.82 | 1 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 9995 | 771 | 0 | 1 | 39 | 5 | 0.00 | 2 | |
| 9996 | 516 | 0 | 1 | 35 | 10 | 57369.61 | 1 | |
| 9997 | 709 | 0 | 0 | 36 | 7 | 0.00 | 1 | |
| 9998 | 772 | 1 | 1 | 42 | 3 | 75075.31 | 2 | |
| 9999 | 792 | 0 | 0 | 28 | 4 | 130142.79 | 1 | |

10000 rows × 11 columns

```
x=df.iloc[:, :-1]  
x
```

| | CreditScore | Geography | Gender | Age | Tenure | Balance | NumOfProducts | HasCrC |
|------|-------------|-----------|--------|-----|--------|-----------|---------------|--------|
| 0 | 619 | 0 | 0 | 42 | 2 | 0.00 | 1 | |
| 1 | 608 | 2 | 0 | 41 | 1 | 83807.86 | 1 | |
| 2 | 502 | 0 | 0 | 42 | 8 | 159660.80 | 3 | |
| 3 | 699 | 0 | 0 | 39 | 1 | 0.00 | 2 | |
| 4 | 850 | 2 | 0 | 43 | 2 | 125510.82 | 1 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 9995 | 771 | 0 | 1 | 39 | 5 | 0.00 | 2 | |
| 9996 | 516 | 0 | 1 | 35 | 10 | 57369.61 | 1 | |
| 9997 | 709 | 0 | 0 | 36 | 7 | 0.00 | 1 | |
| 9998 | 772 | 1 | 1 | 42 | 3 | 75075.31 | 2 | |
| 9999 | 792 | 0 | 0 | 28 | 4 | 130142.79 | 1 | |

10000 rows × 10 columns

```
from sklearn.preprocessing import StandardScaler  
sc=StandardScaler()  
x=sc.fit_transform(x)
```

x

```
array([[ -0.32622142, -0.90188624, -1.09598752, ...,  0.64609167,  
        0.97024255,  0.02188649],  
       [ -0.44003595,  1.51506738, -1.09598752, ..., -1.54776799,  
        0.97024255,  0.21653375],  
       [ -1.53679418, -0.90188624, -1.09598752, ...,  0.64609167,  
       -1.03067011,  0.2406869 ],  
       ...,  
       [  0.60498839, -0.90188624, -1.09598752, ..., -1.54776799,  
        0.97024255, -1.00864308],  
       [  1.25683526,  0.30659057,  0.91241915, ...,  0.64609167,
```

```

-1.03067011, -0.12523071],
[ 1.46377078, -0.90188624, -1.09598752, ...,  0.64609167,
-1.03067011, -1.07636976]])

y=df["Exited"].values
y

array([1, 0, 1, ..., 1, 1, 0])

from sklearn.model_selection import train_test_split
xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.30,random_state=1)

import tensorflow as tf
from tensorflow.keras import Sequential # it is used to build NN
from tensorflow.keras.layers import Dense # it is used to add hidden layers
from sklearn.metrics import classification_report #evaluation

#builind NN

#step 1-: init a model
ann=Sequential()

#step2-: add layers to model
ann.add(Dense(units=12,activation="relu")) #HL
ann.add(Dense(units=1,activation="sigmoid"))

#step3-: establish connection with layers
ann.compile(optimizer="adam",loss="binary_crossentropy",metrics=["accuracy"])

#step4-: fit the model
ann.fit(xtrain,ytrain,epochs=100,batch_size=32)

#step5-: make predictions
ypred=ann.predict(xtest)

Epoch 1/100
219/219 [=====] - 1s 2ms/step - loss: 0.6460 - accuracy: 0.6351
Epoch 2/100
219/219 [=====] - 0s 2ms/step - loss: 0.4658 - accuracy: 0.7999
Epoch 3/100
219/219 [=====] - 0s 2ms/step - loss: 0.4357 - accuracy: 0.8109
Epoch 4/100
219/219 [=====] - 0s 2ms/step - loss: 0.4221 - accuracy: 0.8203
Epoch 5/100
219/219 [=====] - 0s 2ms/step - loss: 0.4114 - accuracy: 0.8274
Epoch 6/100
219/219 [=====] - 0s 2ms/step - loss: 0.4014 - accuracy: 0.8324
Epoch 7/100
219/219 [=====] - 0s 2ms/step - loss: 0.3922 - accuracy: 0.8397
Epoch 8/100
219/219 [=====] - 0s 2ms/step - loss: 0.3839 - accuracy: 0.8457
Epoch 9/100
219/219 [=====] - 0s 2ms/step - loss: 0.3770 - accuracy: 0.8499
Epoch 10/100
219/219 [=====] - 0s 2ms/step - loss: 0.3714 - accuracy: 0.8506
Epoch 11/100
219/219 [=====] - 0s 2ms/step - loss: 0.3672 - accuracy: 0.8519
Epoch 12/100
219/219 [=====] - 0s 2ms/step - loss: 0.3639 - accuracy: 0.8526
Epoch 13/100
219/219 [=====] - 0s 2ms/step - loss: 0.3613 - accuracy: 0.8536
Epoch 14/100
219/219 [=====] - 0s 2ms/step - loss: 0.3597 - accuracy: 0.8551
Epoch 15/100
219/219 [=====] - 0s 2ms/step - loss: 0.3580 - accuracy: 0.8541
Epoch 16/100
219/219 [=====] - 0s 2ms/step - loss: 0.3568 - accuracy: 0.8559
Epoch 17/100
219/219 [=====] - 0s 2ms/step - loss: 0.3559 - accuracy: 0.8546
Epoch 18/100
219/219 [=====] - 0s 2ms/step - loss: 0.3550 - accuracy: 0.8543
Epoch 19/100
219/219 [=====] - 1s 2ms/step - loss: 0.3544 - accuracy: 0.8567
Epoch 20/100
219/219 [=====] - 1s 2ms/step - loss: 0.3537 - accuracy: 0.8561
Epoch 21/100
219/219 [=====] - 1s 2ms/step - loss: 0.3531 - accuracy: 0.8556

```

```
Epoch 22/100
219/219 [=====] - 0s 2ms/step - loss: 0.3527 - accuracy: 0.8556
Epoch 23/100
219/219 [=====] - 1s 2ms/step - loss: 0.3523 - accuracy: 0.8571
Epoch 24/100
219/219 [=====] - 0s 2ms/step - loss: 0.3516 - accuracy: 0.8566
Epoch 25/100
219/219 [=====] - 0s 1ms/step - loss: 0.3509 - accuracy: 0.8576
Epoch 26/100
219/219 [=====] - 0s 2ms/step - loss: 0.3508 - accuracy: 0.8566
Epoch 27/100
219/219 [=====] - 0s 2ms/step - loss: 0.3501 - accuracy: 0.8576
Epoch 28/100
219/219 [=====] - 0s 2ms/step - loss: 0.3497 - accuracy: 0.8560
Epoch 29/100
219/219 [=====] - 0s 2ms/step - loss: 0.3495 - accuracy: 0.8571
```

ypred

```
array([[0.03889292],
       [0.07804814],
       [0.07518935],
       ...,
       [0.03322417],
       [0.06540847],
       [0.03827015]], dtype=float32)
```

```
#step6-: Set the threshold
ypred=np.where(ypred>0.5,1,0)
ypred
```

```
array([[0],
       [0],
       [0],
       ...,
       [0],
       [0],
       [0]])
```

```
print(classification_report(ytest,ypred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.88 | 0.96 | 0.92 | 2373 |
| 1 | 0.76 | 0.48 | 0.59 | 627 |
| accuracy | | | 0.86 | 3000 |
| macro avg | 0.82 | 0.72 | 0.75 | 3000 |
| weighted avg | 0.85 | 0.86 | 0.85 | 3000 |